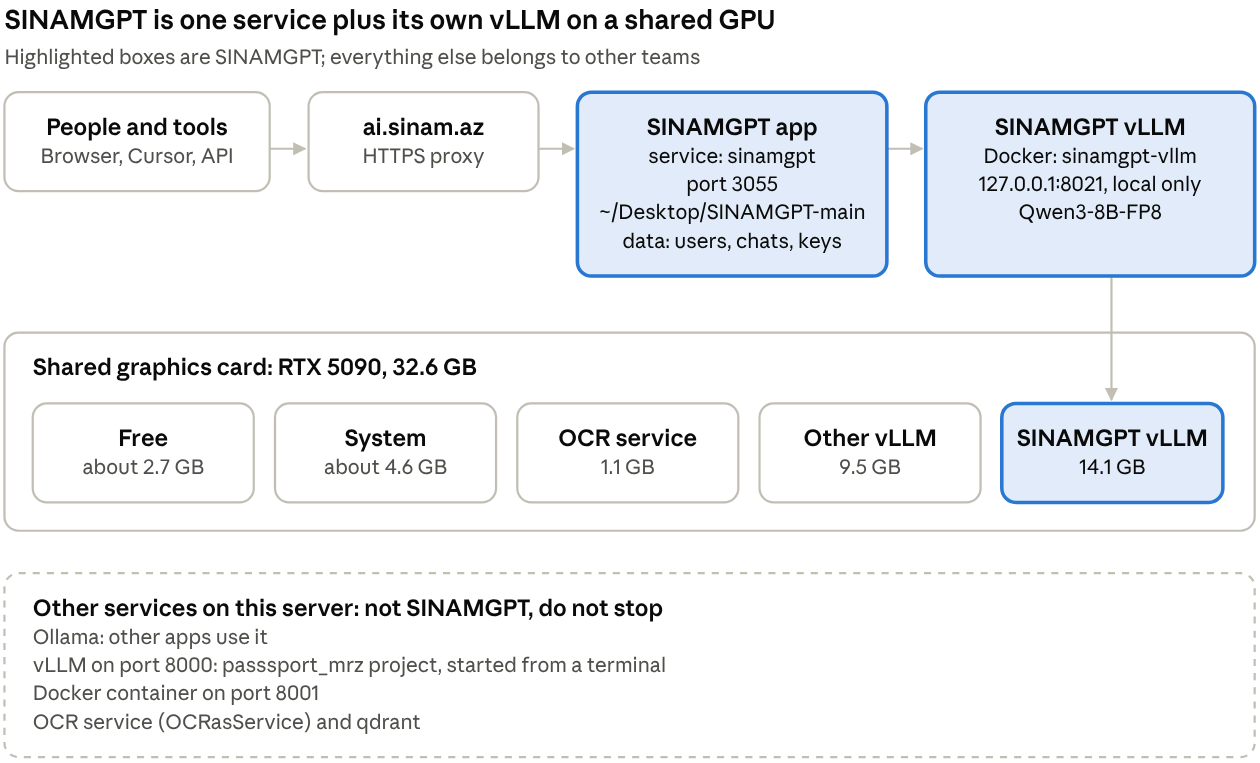
SINAMGPT Server Health Check

Oct 21, 2026 · @Emil Talıbov

Server map



server map · SINAMGPT app, its vLLM and the shared GPU

Requests reach the app through ai.sinam.az; the app talks only to its own vLLM on 127.0.0.1:8021. Graphics card shares as measured on 2026-10-21.

Quick health check

Run this block on the server, in a terminal or over SSH (`ssh aiadmin02@10.10.206.2`). It only reads, changes nothing, and asks for the sudo password once.

```
echo "== 1. SINAMGPT app ==" ; echo "starts at boot: $(systemctl is-enabled
sinamgpt)" ; echo "running: $(systemctl is-active sinamgpt)" ; curl -s -o
/dev/null -w "local: %{http_code}\n" http://127.0.0.1:3055/login ; curl -s -o
/dev/null -w "domain: %{http_code}\n" https://ai.sinam.az/login
echo "== 2. vLLM ==" ; sudo docker ps --filter name=sinamgpt-vllm --format
'{{.Status}}  {{.Ports}}' ; sudo docker inspect -f 'restarts by itself:
{{.HostConfig.RestartPolicy.Name}}' sinamgpt-vllm-vllm-1 ; curl -s -o
/dev/null -w "answers with key: %{http_code}\n" -H "Authorization: Bearer
$(grep '^VLLM_API_KEY=' ~/Desktop/sinamgpt-vllm/.env.vllm | cut -d= -f2)"
http://127.0.0.1:8021/v1/models
echo "== 3. Docker starts at boot ==" ; systemctl is-enabled docker
echo "== 4. Graphics card ==" ; nvidia-smi --query-gpu=memory.used,memory.free
--format=csv,noheader
echo "== 5. Memory ==" ; free -h | head -3
echo "== 6. Disk ==" ; df -h / | tail -1
echo "== 7. Ollama for other apps ==" ; systemctl is-active ollama
echo "== 8. App errors in the last hour ==" ; sudo journalctl -u sinamgpt --
since "-1h" --no-pager | grep -ciE 'error|unhandled|failed'
```

What good results look like

All eight checks matched these values on 2026-10-21; a row that differs points to a fix in the next section.

#	Check	Good result
1	App	`enabled`, `active`, `local: 200`, `domain: 200`
2	vLLM	`Up … (healthy) 127.0.0.1:8021->8000/tcp`, `unless-stopped`, `answers with key: 200`
3	Docker at boot	`enabled`, so vLLM comes back after a reboot
4	Graphics card	about 29,350 MiB used, about 2,760 MiB free
5	Memory	30 GB or more available; a full swap is harmless while that much RAM is free
6	Disk	under 80% used (29% on 2026-10-21)
7	Ollama	`active`, because other apps use it
8	App errors in the last hour	`0` or a few

If something is wrong

Start with the matching fix below, then run the quick check again. Never restart the whole server or Docker first: other teams run services here.

App is down (check 1)

1. Restart it: `sudo systemctl restart sinamgpt`
2. Read why it stopped: `sudo journalctl -u sinamgpt -n 50 --no-pager`
3. Never start a second copy by hand (`npm run start` or `nohup`) while the service runs: both want port 3055.

vLLM is down or not healthy (check 2)

1. Start it: `cd ~/Desktop/sinamgpt-vllm && sudo docker compose -p sinamgpt-vllm --env-file .env.vllm -f docker-compose.vllm.yml up -d`
2. Wait 1–5 minutes until the status says `(healthy)`. The log: same command with `logs --tail=60 vllm` instead of `up -d`.
3. Log says free memory is less than desired: another program took graphics memory. Wait until `ollama ps` lists no model, then start again.

4. Log mentions a GPU vendor or CDI: run `sudo nvidia-ctk cdi generate --output=/etc/cdi/nvidia.yaml`. Needed again after every NVIDIA driver update.

Graphics card almost full (check 4 shows under 1,000 MiB free)

1. See who holds it: `nvidia-smi --query-compute-apps=pid,process_name,used_memory --format=csv`
2. Normal shares: SINAMGPT vLLM about 14.1 GB, the other team's vLLM about 9.5 GB, the OCR service about 1.1 GB.
3. Do not stop other teams' programs. If SINAMGPT needs to give room, lower `VLLM_GPU_MEMORY_UTILIZATION` in `~/Desktop/sinamgpt-vllm/.env.vllm` (0.42 today) and restart vLLM.

Memory low (check 5 shows under 5 GB available)

1. See the biggest users: `ps -eo pid,user,%mem,rss,etime,comm --sort=-rss | head -12`
2. The elementary App Center (`io.elementary.appcenter`) grew to about 2 GB per copy on 2026-10-21. Closing it is safe: `sudo kill <pid>`.

Remote desktop frozen or slow

1. Use SSH from your laptop instead: `ssh aiadmin02@10.10.206.2`. It keeps working when the desktop does not.
2. Check memory as above; a long-running desktop session pushes itself into swap.
3. Do not log out of that desktop while the vLLM on port 8000 runs from a terminal inside it: logging out may stop it.

Everyday commands

Use these instead of starting or killing processes by hand; the vLLM commands run from `~/Desktop/sinamgpt-vllm`.

Task	Command
App status	`systemctl status sinamgpt --no-pager`
Restart the app	`sudo systemctl restart sinamgpt`

Task	Command
App log	`sudo journalctl -u sinamgpt -n 50 --no-pager`
vLLM status	`sudo docker ps --filter name=sinamgpt-vllm`
Restart vLLM	`sudo docker compose -p sinamgpt-vllm --env-file .env.vllm -f docker-compose.vllm.yml restart`
vLLM log	`sudo docker compose -p sinamgpt-vllm --env-file .env.vllm -f docker-compose.vllm.yml logs --tail=60 vllm`

Update the app to a new version

1. Copy the new files into `~/Desktop/SINAMGPT-main`, but never `data/`, `.env.local`, `node_modules` or `.next`.
2. Stop the app: `sudo systemctl stop sinamgpt`
3. Back up data and settings: `BK=~/sinamgpt-backup-$(date +%Y%m%d-%H%M) && mkdir -p "$BK" && cp -a ~/Desktop/SINAMGPT-main/data ~/Desktop/SINAMGPT-main/.env.local "$BK"/`
4. Install and build: `cd ~/Desktop/SINAMGPT-main && npm ci --include=dev --ignore-scripts && npm run build`
5. Start it again: `sudo systemctl start sinamgpt`, then run the quick check.

Rules that keep it safe

- Do not stop Ollama, the vLLM on port 8000, the container on port 8001 or the OCR service: they belong to other apps.
- Never run `docker compose down -v` for SINAMGPT's vLLM: `-v` deletes the downloaded model.
- Files copied from Windows carry hidden line-end characters; clean a config file before use with `sed -i 's/\r$//' <file>`.
- Never paste an API key into chat or a ticket. A leaked key: Developer, then API keys, then Revoke.